Autoregressive rollout updates both levels of a fixed-size context state

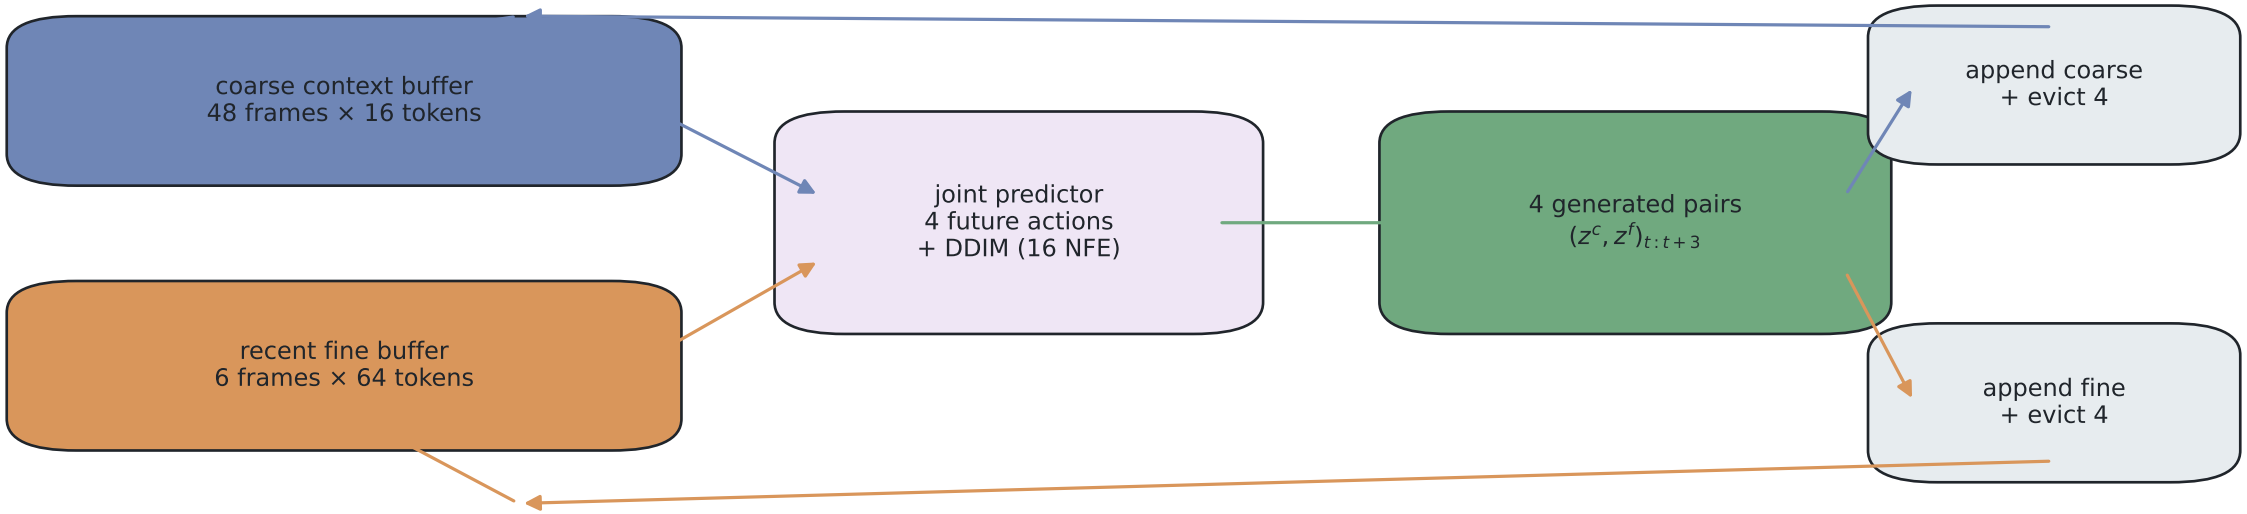

advance by four frames; both buffers retain the same fixed geometry at every step

context state = 48×256 + 6×1,024 = 18,432 latent numbers, independent of rollout length

48 ground-truth warm-up frames	152 self-fed frames, generated four at a time

The layout matches training, but self-fed values need not follow the clean-context training distribution.