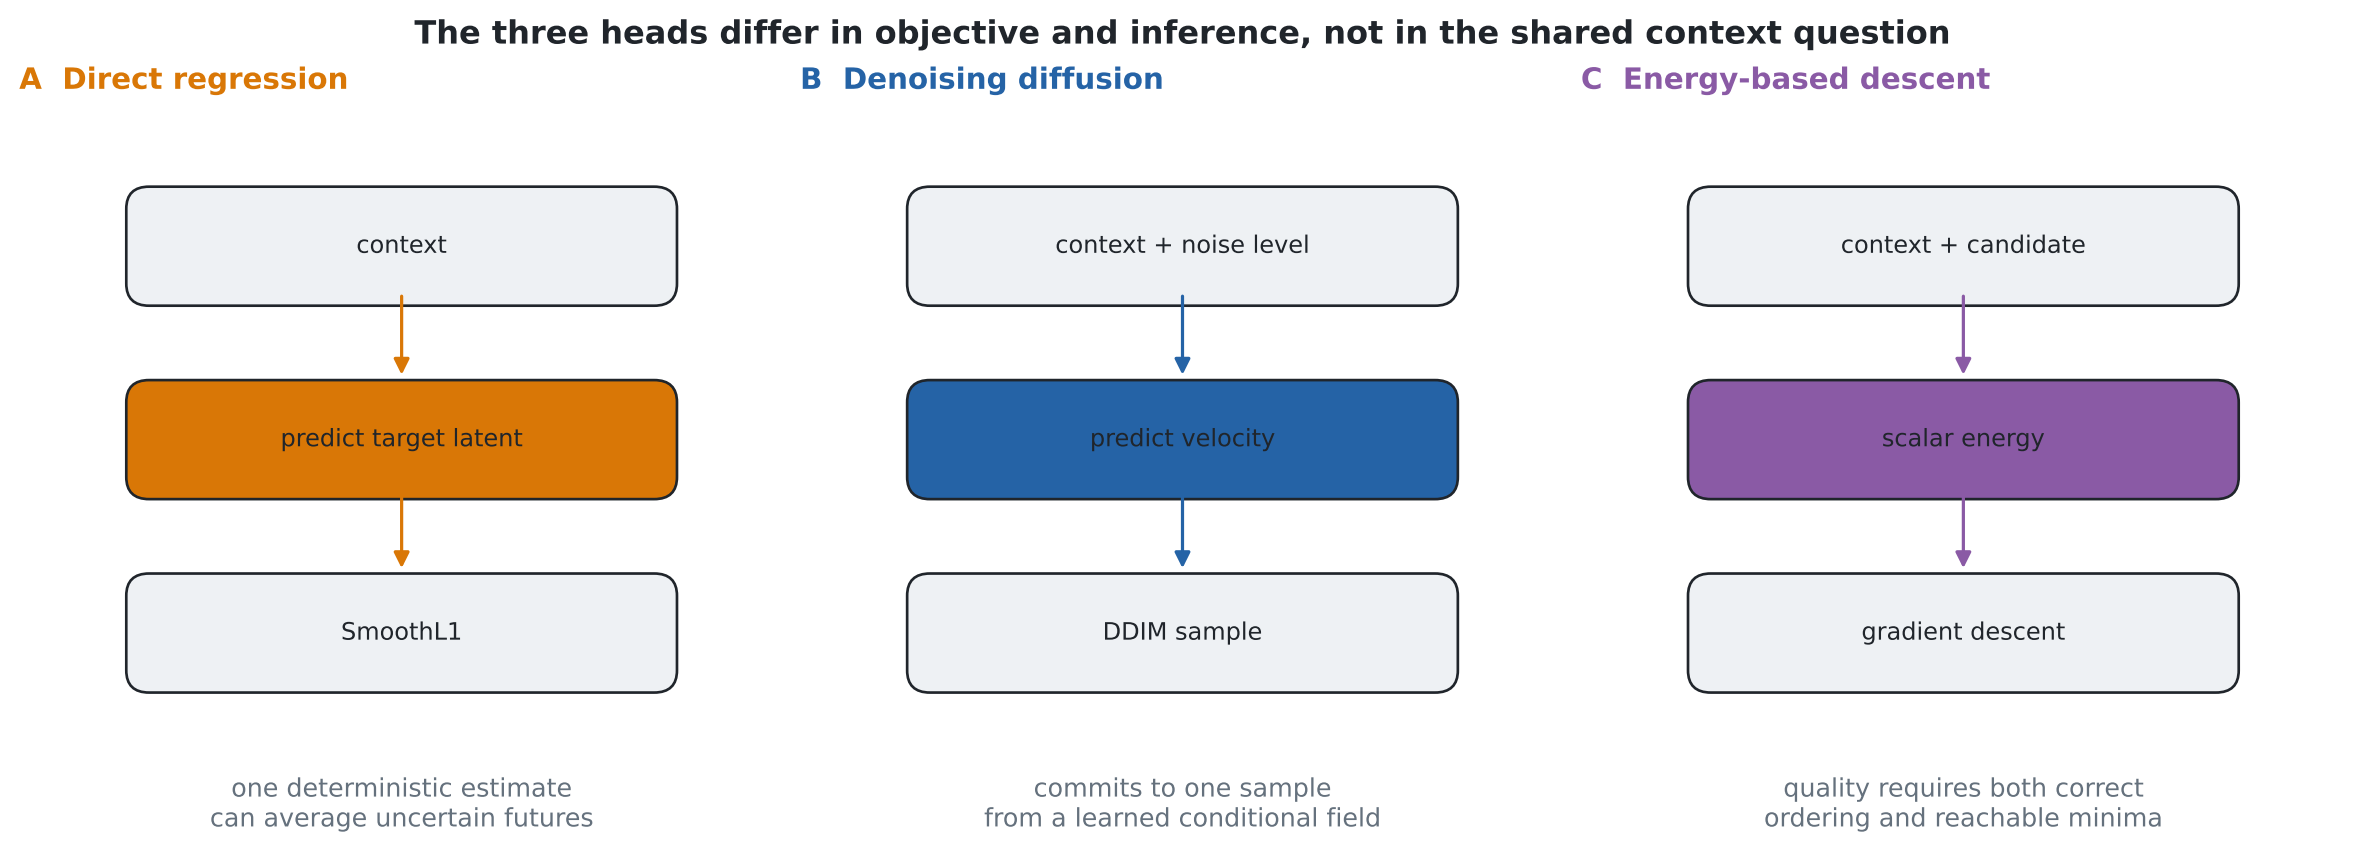
The three heads differ in objective and inference, not in the shared context question
A Direct regression
context
predict target latent
SmoothL1
one deterministic estimate
can average uncertain futures
B Denoising diffusion
context + noise level
predict velocity
DDIM sample
commits to one sample
from a learned conditional field
C Energy-based descent
context + candidate
scalar energy
gradient descent
quality requires both correct
ordering and reachable minima