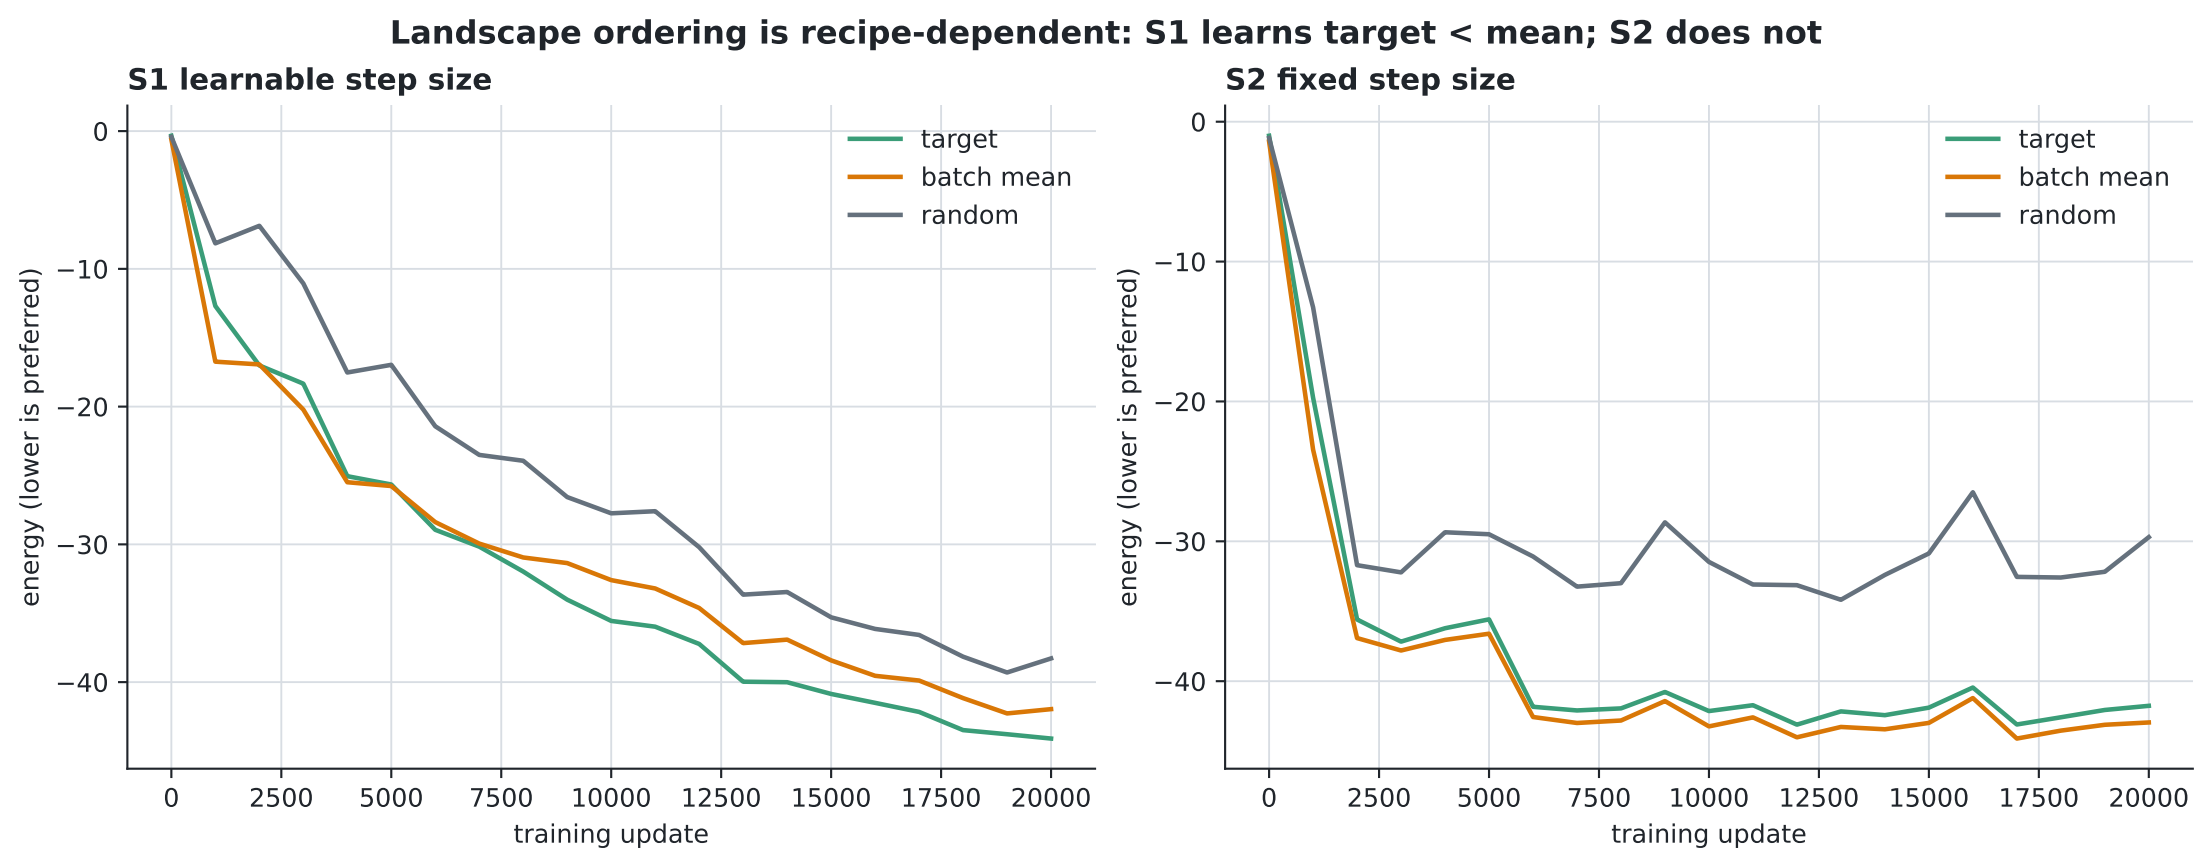
Landscape ordering is recipe-dependent: S1 learns target < mean; S2 does not
S1 learnable step size
S2 fixed step size
target
batch mean
random
energy (lower is preferred)
training update
0
−10
−20
−30
−40
0
2500
5000
7500
10000
12500
15000
17500
20000